

# Bayesian Deep Learning and Uncertainty

Christos Louizos, Max Welling



University of Amsterdam

# Overview

- Part 1:
  - Structured approximate posteriors in Bayesian neural nets
- Part 2:
  - Exploring model uncertainty

## **Part 1**

# Structured and Efficient Variational Deep Learning with Matrix Gaussian Posteriors

ICML 2016

# VI for Bayesian neural networks (BNNs)

# VI for Bayesian neural networks (BNNs)

Stochastic gradient VI for BNNs optimizes:

$$\mathbb{E}_{\prod_{i=1}^H q(\mathbf{W}_i)} [\log p(\mathbf{y}|\mathbf{x}, \mathbf{W}_{1:H})] - \sum_{i=1}^H KL(q(\mathbf{W}_i), p(\mathbf{W}_i))$$

w.r.t. to the posterior distributions  $q$

# VI for Bayesian neural networks (BNNs)

Stochastic gradient VI for BNNs optimizes:

$$\mathbb{E}_{\prod_{i=1}^H q(\mathbf{W}_i)} [\log p(\mathbf{y}|\mathbf{x}, \mathbf{W}_{1:H})] - \sum_{i=1}^H KL(q(\mathbf{W}_i), p(\mathbf{W}_i))$$

w.r.t. to the posterior distributions  $q$

Usual choices for  $q$  are fully factorized Gaussians

$$q(\mathbf{W}_i) = \prod_{r=1}^R \prod_{c=1}^C \mathcal{N}(\mu_{rc}, \sigma_{rc}^2)$$

# Implications of fully factorized posteriors

# Implications of fully factorized posteriors

- Doubles the amount of trainable parameters
  - Prohibitive for large neural networks



# Implications of fully factorized posteriors


- Doubles the amount of trainable parameters
  - Prohibitive for large neural networks
- Too much “noise” due to a lot of independent sources of variation

# Implications of fully factorized posteriors

- Doubles the amount of trainable parameters
  - Prohibitive for large neural networks
- Too much “noise” due to a lot of independent sources of variation
- No information sharing across weights

# Implications of fully factorized posteriors

- Doubles
  - Prohibi
- Too muc  
variation
- No infor



urces of

Fully factorized VI BNN

# Structured VI for BNNs

How we can introduce structure efficiently?

# Structured VI for BNNs

How we can introduce structure efficiently?

## **Simple idea**

Instead of approximating the marginal of each element in the weight matrix approximate the weight matrix directly!

# Structured VI for BNNs

How we can introduce structure efficiently?

## Simple idea

Instead of approximating the marginal of each element in the weight matrix approximate the weight matrix directly!

Assume a matrix Gaussian posterior over each weight matrix:

$$\begin{aligned} q(\mathbf{W}_i) &= \mathcal{MN}(\mathbf{M}_{r \times c}, \mathbf{U}_{r \times r}, \mathbf{V}_{c \times c}) \\ &= \mathcal{N}(\text{vec}(\mathbf{M}), \mathbf{U} \otimes \mathbf{V}) \end{aligned}$$

$\mathbf{M}$  corresponds to the mean matrix

$\mathbf{U}$  corresponds to the row covariance

$\mathbf{V}$  corresponds to the column covariance

# Structured VI with matrix Gaussian posteriors

# Structured VI with matrix Gaussian posteriors

- Compactly model correlations among weights



# Structured VI with matrix Gaussian posteriors


- Compactly model correlations among weights
- “Full” matrix Gaussian is still expensive

# Structured VI with matrix Gaussian posteriors

- Compactly model correlations among weights
- “Full” matrix Gaussian is still expensive
- Low rank approximations for **U**, **V**
  - Diagonal approximations reduce parameters greatly
  - Still maintain some correlations among weights

# Structured VI with matrix Gaussian posteriors

- Comparison
- “Full” model
- Low rank
- Diagonal
- Still more

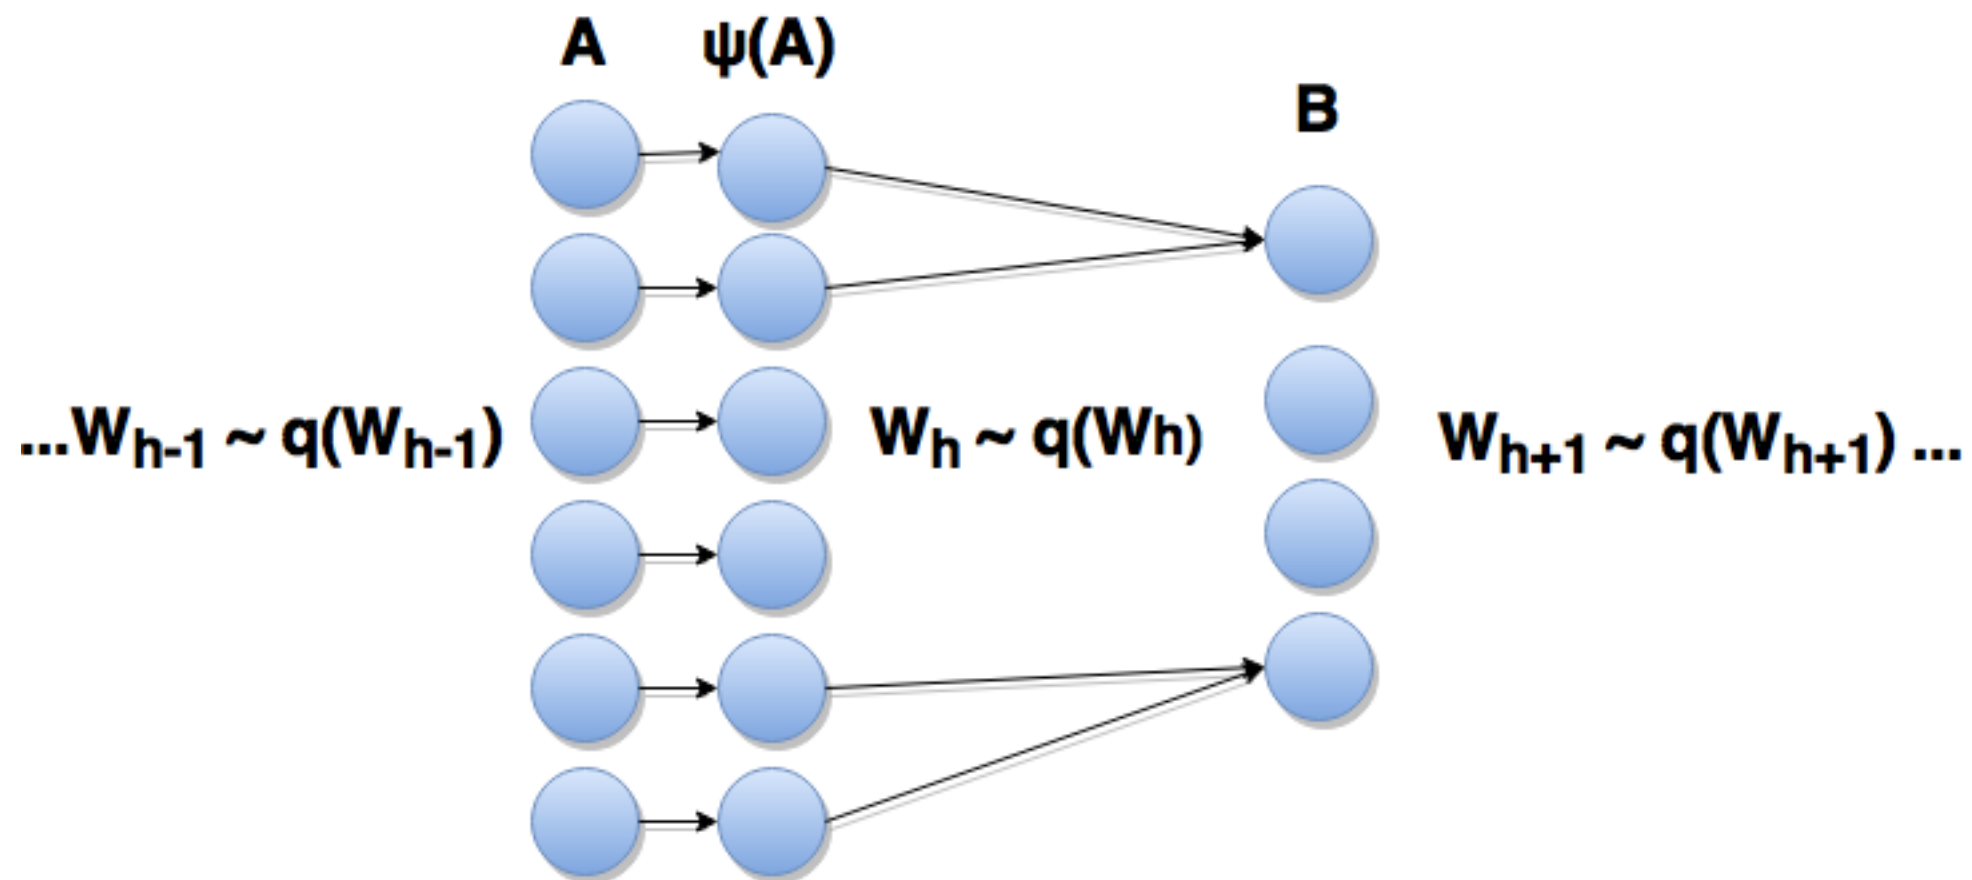


Structured VI BNN

# Global posterior sampling


# Global posterior sampling

In order to perform stochastic gradient VI we have to sample the approximate weight posterior at each layer:



# Global posterior sampling

In order to perform stochastic gradient VI we have to sample the approximate weight posterior at each layer:



where **A** is a mini batch of the previous layer output,  $\psi(\cdot)$  is an element wise nonlinearity and **B** is the layer output

# Local posterior sampling

# Local posterior sampling

[1] showed that we can obtain a lower variance gradient estimator if we instead sample the pre-activation variables **B** by locally marginalising the weights, i.e:



# Local posterior sampling

[1] showed that we can obtain a lower variance gradient estimator if we instead sample the pre-activation variables  $\mathbf{B}$  by locally marginalising the weights, i.e:

$$\mathbf{W}_h \sim \prod_{r=1}^R \prod_{c=1}^C \mathcal{N}(\mu_{rc}, \sigma_{rc}^2)$$
$$\mathbf{B} = \psi(\mathbf{A})\mathbf{W}_h$$

# Local posterior sampling

[1] showed that we can obtain a lower variance gradient estimator if we instead sample the pre-activation variables  $\mathbf{B}$  by locally marginalising the weights, i.e:

$$\mathbf{W}_h \sim \prod_{r=1}^R \prod_{c=1}^C \mathcal{N}(\mu_{rc}, \sigma_{rc}^2)$$
$$\mathbf{B} = \psi(\mathbf{A})\mathbf{W}_h$$

is equivalent to:

# Local posterior sampling

[1] showed that we can obtain a lower variance gradient estimator if we instead sample the pre-activation variables  $\mathbf{B}$  by locally marginalising the weights, i.e:

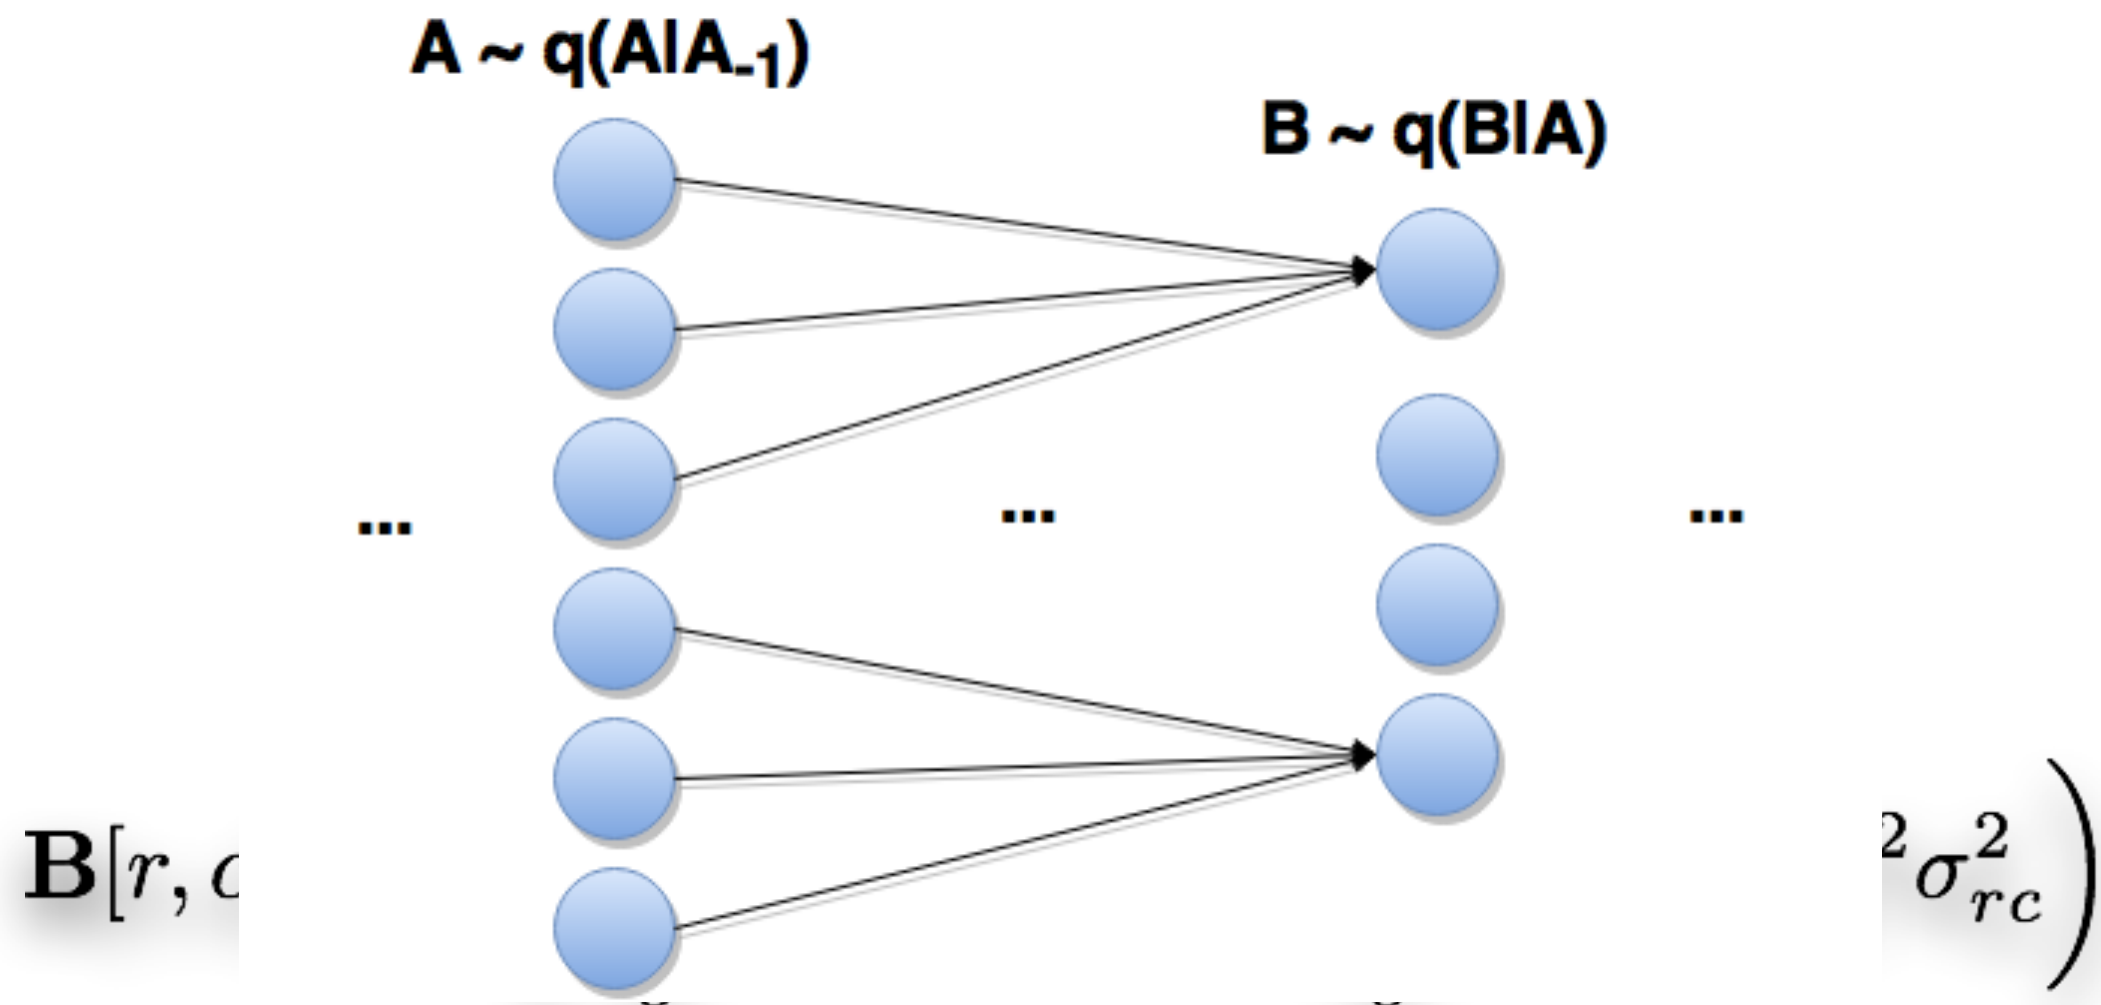
$$\mathbf{W}_h \sim \prod_{r=1}^R \prod_{c=1}^C \mathcal{N}(\mu_{rc}, \sigma_{rc}^2)$$
$$\mathbf{B} = \psi(\mathbf{A})\mathbf{W}_h$$

is equivalent to:

$$\mathbf{B}[r, c] \sim \mathcal{N}\left(\sum_c \psi(\mathbf{A})[r, c] \mu_{rc}, \sum_c \psi(\mathbf{A})[r, c]^2 \sigma_{rc}^2\right)$$

# Local posterior sampling

[1] showed that we can obtain a lower variance gradient estimator if we instead sample the pre-activation variables **B** by locally marginalising the weights, i.e:



# Local posterior sampling with a Matrix Gaussian

# Local posterior sampling with a Matrix Gaussian

Similarly we have that:

$$\mathbf{W}_h \sim \mathcal{MN}(\mathbf{M}_h, \mathbf{U}_h, \mathbf{V}_h)$$

$$\mathbf{B} = \psi(\mathbf{A})\mathbf{W}_h$$

# Local posterior sampling with a Matrix Gaussian

Similarly we have that:

$$\mathbf{W}_h \sim \mathcal{MN}(\mathbf{M}_h, \mathbf{U}_h, \mathbf{V}_h)$$

$$\mathbf{B} = \psi(\mathbf{A})\mathbf{W}_h$$

is equivalent to:

$$\mathbf{B} \sim \mathcal{MN}(\psi(\mathbf{A})\mathbf{M}_h, \psi(\mathbf{A})\mathbf{U}_h\psi(\mathbf{A})^T, \mathbf{V}_h)$$

# Local posterior sampling with a Matrix Gaussian

Similarly we have that:

$$\mathbf{W}_h \sim \mathcal{MN}(\mathbf{M}_h, \mathbf{U}_h, \mathbf{V}_h)$$

$$\mathbf{B} = \psi(\mathbf{A})\mathbf{W}_h$$

is equivalent to:

$$\mathbf{B} \sim \mathcal{MN}(\psi(\mathbf{A})\mathbf{M}_h, \psi(\mathbf{A})\mathbf{U}_h\psi(\mathbf{A})^T, \mathbf{V}_h)$$

or else:

$$\mathbf{B} \sim \mathcal{N}(\text{vec}(\psi(\mathbf{A})\mathbf{M}_h), \mathbf{K}(\mathbf{A}, \mathbf{A}))$$

$$\mathbf{K}(\mathbf{A}, \mathbf{A}) = \mathbf{K}_{out} \otimes \mathbf{K}_{in}(\mathbf{A}, \mathbf{A})$$



# Local posterior sampling with a Matrix Gaussian

Similarly we have that:

$$\mathbf{W}_h \sim \mathcal{MN}(\mathbf{M}_h, \mathbf{U}_h, \mathbf{V}_h)$$

$$\mathbf{B} = \psi(\mathbf{A})\mathbf{W}_h$$

is equivalent to:

$$\mathbf{B} \sim \mathcal{MN}(\psi(\mathbf{A})\mathbf{M}_h, \psi(\mathbf{A})\mathbf{U}_h\psi(\mathbf{A})^T, \mathbf{V}_h)$$

or else:

$$\mathbf{B} \sim \mathcal{N}(\text{vec}(\psi(\mathbf{A})\mathbf{M}_h), \mathbf{K}(\mathbf{A}, \mathbf{A}))$$

$$\mathbf{K}(\mathbf{A}, \mathbf{A}) = \mathbf{K}_{out} \otimes \mathbf{K}_{in}(\mathbf{A}, \mathbf{A})$$

where:

$$\mathbf{K}_{in}(\mathbf{A}, \mathbf{A}) = \psi(\mathbf{A})\mathbf{U}_h\psi(\mathbf{A})^T$$

$$\mathbf{K}_{out} = \mathbf{V}_h$$

# Local posterior sampling with a Matrix Gaussian

Similarly we have that:

$$\mathbf{B} \sim \mathcal{N}(\text{vec}(\psi(\mathbf{A})\mathbf{M}_h), \mathbf{K}(\mathbf{A}, \mathbf{A}))$$

So we can view each layer as a finite rank multi-output Gaussian Process and the whole network as a deep Gaussian Process[2] with a specific kernel!

or else:

$$\mathbf{B} \sim \mathcal{N}(\text{vec}(\psi(\mathbf{A})\mathbf{M}_h), \mathbf{K}(\mathbf{A}, \mathbf{A}))$$

$$\mathbf{K}(\mathbf{A}, \mathbf{A}) = \mathbf{K}_{out} \otimes \mathbf{K}_{in}(\mathbf{A}, \mathbf{A})$$

where:

$$\mathbf{K}_{in}(\mathbf{A}, \mathbf{A}) = \psi(\mathbf{A})\mathbf{U}_h\psi(\mathbf{A})^T$$

$$\mathbf{K}_{out} = \mathbf{V}_h$$

# Sampling a matrix Gaussian

# Sampling a matrix Gaussian

In order to sample the matrix Gaussian at each layer we have to do:

$$\mathbf{B} = \psi(\mathbf{A})\mathbf{M}_h + \mathbf{K}_{in}(\mathbf{A}, \mathbf{A})^{\frac{1}{2}} \mathbf{E} \mathbf{K}_{out}^{\frac{1}{2}}$$
$$\mathbf{E}[i, j] \sim \mathcal{N}(0, 1)$$

# Sampling a matrix Gaussian

In order to sample the matrix Gaussian at each layer we have to do:

$$\mathbf{B} = \psi(\mathbf{A})\mathbf{M}_h + \mathbf{K}_{in}(\mathbf{A}, \mathbf{A})^{\frac{1}{2}} \mathbf{E} \mathbf{K}_{out}^{\frac{1}{2}}$$
$$\mathbf{E}[i, j] \sim \mathcal{N}(0, 1)$$

Expensive since we have to compute the square root of the input kernel for each mini-batch.

# Inducing points for BNNs

# Inducing points for BNNs

To scale it up we introduce pseudo input-output pairs and assume conditional independence for the elements of  $\mathbf{B}$  (FITC GP approximation):

$$\mathbf{b} \sim p(\mathbf{b}|\mathbf{a}, \tilde{\mathbf{B}}, \tilde{\mathbf{A}})$$

$$\sim \mathcal{N}(\psi(\mathbf{a})\mathbf{M}_h + \mathbf{K}_{in}(\mathbf{a}, \tilde{\mathbf{A}})\mathbf{K}_{in}(\tilde{\mathbf{A}}, \tilde{\mathbf{A}})^{-1}(\tilde{\mathbf{B}} - \tilde{\mathbf{A}}\mathbf{M}_h),$$

$$\mathbf{K}_{out} \odot (\mathbf{K}_{in}(\mathbf{a}, \mathbf{a}) - \mathbf{K}_{in}(\mathbf{a}, \tilde{\mathbf{A}})\mathbf{K}_{in}(\tilde{\mathbf{A}}, \tilde{\mathbf{A}})^{-1}\mathbf{K}_{in}(\mathbf{a}, \tilde{\mathbf{A}})^T))$$

# Inducing points for BNNs

To scale it up we introduce pseudo input-output pairs and assume conditional independence for the elements of **B** (FITC GP approximation):

$$\mathbf{b} \sim p(\mathbf{b}|\mathbf{a}, \tilde{\mathbf{B}}, \tilde{\mathbf{A}})$$

$$\sim \mathcal{N}(\psi(\mathbf{a})\mathbf{M}_h + \mathbf{K}_{in}(\mathbf{a}, \tilde{\mathbf{A}})\mathbf{K}_{in}(\tilde{\mathbf{A}}, \tilde{\mathbf{A}})^{-1}(\tilde{\mathbf{B}} - \tilde{\mathbf{A}}\mathbf{M}_h),$$

$$\mathbf{K}_{out} \odot (\mathbf{K}_{in}(\mathbf{a}, \mathbf{a}) - \mathbf{K}_{in}(\mathbf{a}, \tilde{\mathbf{A}})\mathbf{K}_{in}(\tilde{\mathbf{A}}, \tilde{\mathbf{A}})^{-1}\mathbf{K}_{in}(\mathbf{a}, \tilde{\mathbf{A}})^T))$$

## Advantages

- Sampling scales cubically with the amount of inducing points
- Variance reduction if real inputs are “similar”



# Regression experiments

# Regression experiments

Dataset	Avg. Test RMSE and Std. Errors				Avg. Test LL and Std. Errors			
	VI	PBP	Dropout	VMG	VI	PBP	Dropout	VMG
Boston	4.32±0.29	3.01±0.18	2.97±0.85	<b>2.81±0.11</b>	-2.90±0.07	-2.57±0.09	<b>-2.46±0.25</b>	-2.54±0.08
Concrete	7.19±0.12	5.67±0.09	5.23±0.53	<b>4.70±0.14</b>	-3.39±0.02	-3.16±0.02	-3.04±0.09	<b>-2.98±0.03</b>
Energy	2.65±0.08	1.80±0.05	1.66±0.19	<b>1.16±0.03</b>	-2.39±0.03	-2.04±0.02	-1.99±0.09	<b>-1.45±0.03</b>
Kin8nm	0.10±0.00	0.10±0.00	0.10±0.00	<b>0.08±0.00</b>	0.90±0.01	0.90±0.01	0.95±0.03	<b>1.14±0.01</b>
Naval	0.01±0.00	0.01±0.00	0.01±0.00	<b>0.00±0.00</b>	3.73±0.12	3.73±0.01	3.80±0.05	<b>5.84±0.00</b>
Pow. Plant	4.33±0.04	4.12±0.03	4.02±0.18	<b>3.88±0.03</b>	-2.89±0.01	-2.84±0.01	-2.80±0.05	<b>-2.78±0.01</b>
Protein	4.84±0.03	4.73±0.01	4.36±0.04	<b>4.14±0.01</b>	-2.99±0.01	-2.97±0.00	-2.89±0.01	<b>-2.84±0.00</b>
Wine	0.65±0.01	0.64±0.01	0.62±0.04	<b>0.61±0.01</b>	-0.98±0.01	-0.97±0.01	-0.93±0.06	<b>-0.93±0.02</b>
Yacht	6.89±0.67	1.02±0.05	1.11±0.38	<b>0.77±0.06</b>	-3.43±0.16	-1.63±0.02	-1.55±0.12	<b>-1.29±0.02</b>
Year	9.034±NA	8.879±NA	8.849±NA	<b>8.780±NA</b>	-3.622±NA	-3.603±NA	<b>-3.588±NA</b>	-3.589±NA

Table 1. Average test set RMSE, predictive log-likelihood and standard errors for the regression datasets. VI, PBP and Dropout correspond to the variational inference method of (Graves, 2011), probabilistic backpropagation (Hernández-Lobato & Adams, 2015) and dropout uncertainty (Gal & Ghahramani, 2015). VMG (Variational Matrix Gaussian) corresponds to the proposed model.

# Classification experiments

# Classification experiments

Method	# layers	Test err.
Max. Likel. (Simard et al., 2003)	$2 \times 800$	1.60
Dropout (Srivastava, 2013)	-	1.25
DropConnect (Wan et al., 2013)	$2 \times 800$	1.20
Bayes B. SM (Blundell et al., 2015)	$2 \times 400$	1.36
	$2 \times 800$	1.34
	$2 \times 1200$	1.32
Var. Dropout (Kingma et al., 2015)	$3 \times 150$	$\approx 1.42$
	$3 \times 250$	$\approx 1.28$
	$3 \times 500$	$\approx 1.18$
	$3 \times 750$	$\approx 1.09$
VMG	$2 \times 400$	<b>1.15</b>
	$3 \times 150$	1.18
	$3 \times 250$	1.11
	$3 \times 500$	1.08
	$3 \times 750$	<b>1.05</b>

Table 2. Test errors for the permutation invariant MNIST dataset. Bayes B. SM correspond to Bayes by Backprop with the scale mixture prior and the variational dropout results are from the Variational (A) model that doesn't downscale the KL-divergence (so as to keep the comparison fair).

## **Part 2**

Model uncertainty from  
BNNs in classification

Opinion: Why are we Bayesian?

# Opinion: Why are we Bayesian?

- Combatting overfitting should not be the primary reason for a Bayesian model!

# Opinion: Why are we Bayesian?

- Combatting overfitting should not be the primary reason for a Bayesian model!
- Being Bayesian is all about handling the uncertainty



# Opinion: Why are we Bayesian?

- Combatting overfitting should not be the primary reason for a Bayesian model!
- Being Bayesian is all about handling the uncertainty
  - We can translate parameter uncertainty into *decision* uncertainty

# Opinion: Why are we Bayesian?

- Combatting overfitting should not be the primary reason for a Bayesian model!
- Being Bayesian is all about handling the uncertainty
  - We can translate parameter uncertainty into *decision* uncertainty
- Model uncertainty is crucial for applications that involve decisions

# Opinion: Why are we Bayesian?

- Combatting overfitting should not be the primary reason for a Bayesian model!
- Being Bayesian is all about handling the uncertainty
  - We can translate parameter uncertainty into *decision* uncertainty
- Model uncertainty is crucial for applications that involve decisions
  - Self-driving cars

# Opinion: Why are we Bayesian?

- Combatting overfitting should not be the primary reason for a Bayesian model!
- Being Bayesian is all about handling the uncertainty
  - We can translate parameter uncertainty into *decision* uncertainty
- Model uncertainty is crucial for applications that involve decisions
  - Self-driving cars
  - Medical applications

# How well do deep BNNs capture uncertainty?


- A simple way to measure the quality
  - Visualize class distribution on perturbed unseen inputs[3]
    - Rotations, translations, scalings
  - Ideally the model should be uncertain
- Experiment with Weight decay, Dropout, Matrix Gaussian posteriors, SGLD, Bootstrap

# Uncertainty testbed

- Simple LeNet network on MNIST
  - Two convolutional, two fully connected layers
- Dropout network: 0.5 dropout rate for all layers[4]
- Matrix Gaussian network
  - 20 pseudo patches/inputs for convolutional, fully connected layers respectively
- 100 samples for SGLD and 20 for Bootstrap
- All networks (except Bootstrap) had  $N(0, 1)$  priors over the parameters


# Weight decay rotations

# Weight decay rotations






# Weight decay rotations




# Dropout rotations

# Dropout rotations




# Dropout rotations




# Matrix Gaussian rotations

# Matrix Gaussian rotations




# Matrix Gaussian rotations




# SGLD rotations



# SGLD rotations




# SGLD rotations



# Bootstrap rotations

# Bootstrap rotations




# Bootstrap rotations



# Weight decay translations

# Weight decay translations




# Weight decay translations






# Dropout translations

# Dropout translations




# Dropout translations




# Matrix Gaussian translations

# Matrix Gaussian translations




# Matrix Gaussian translations



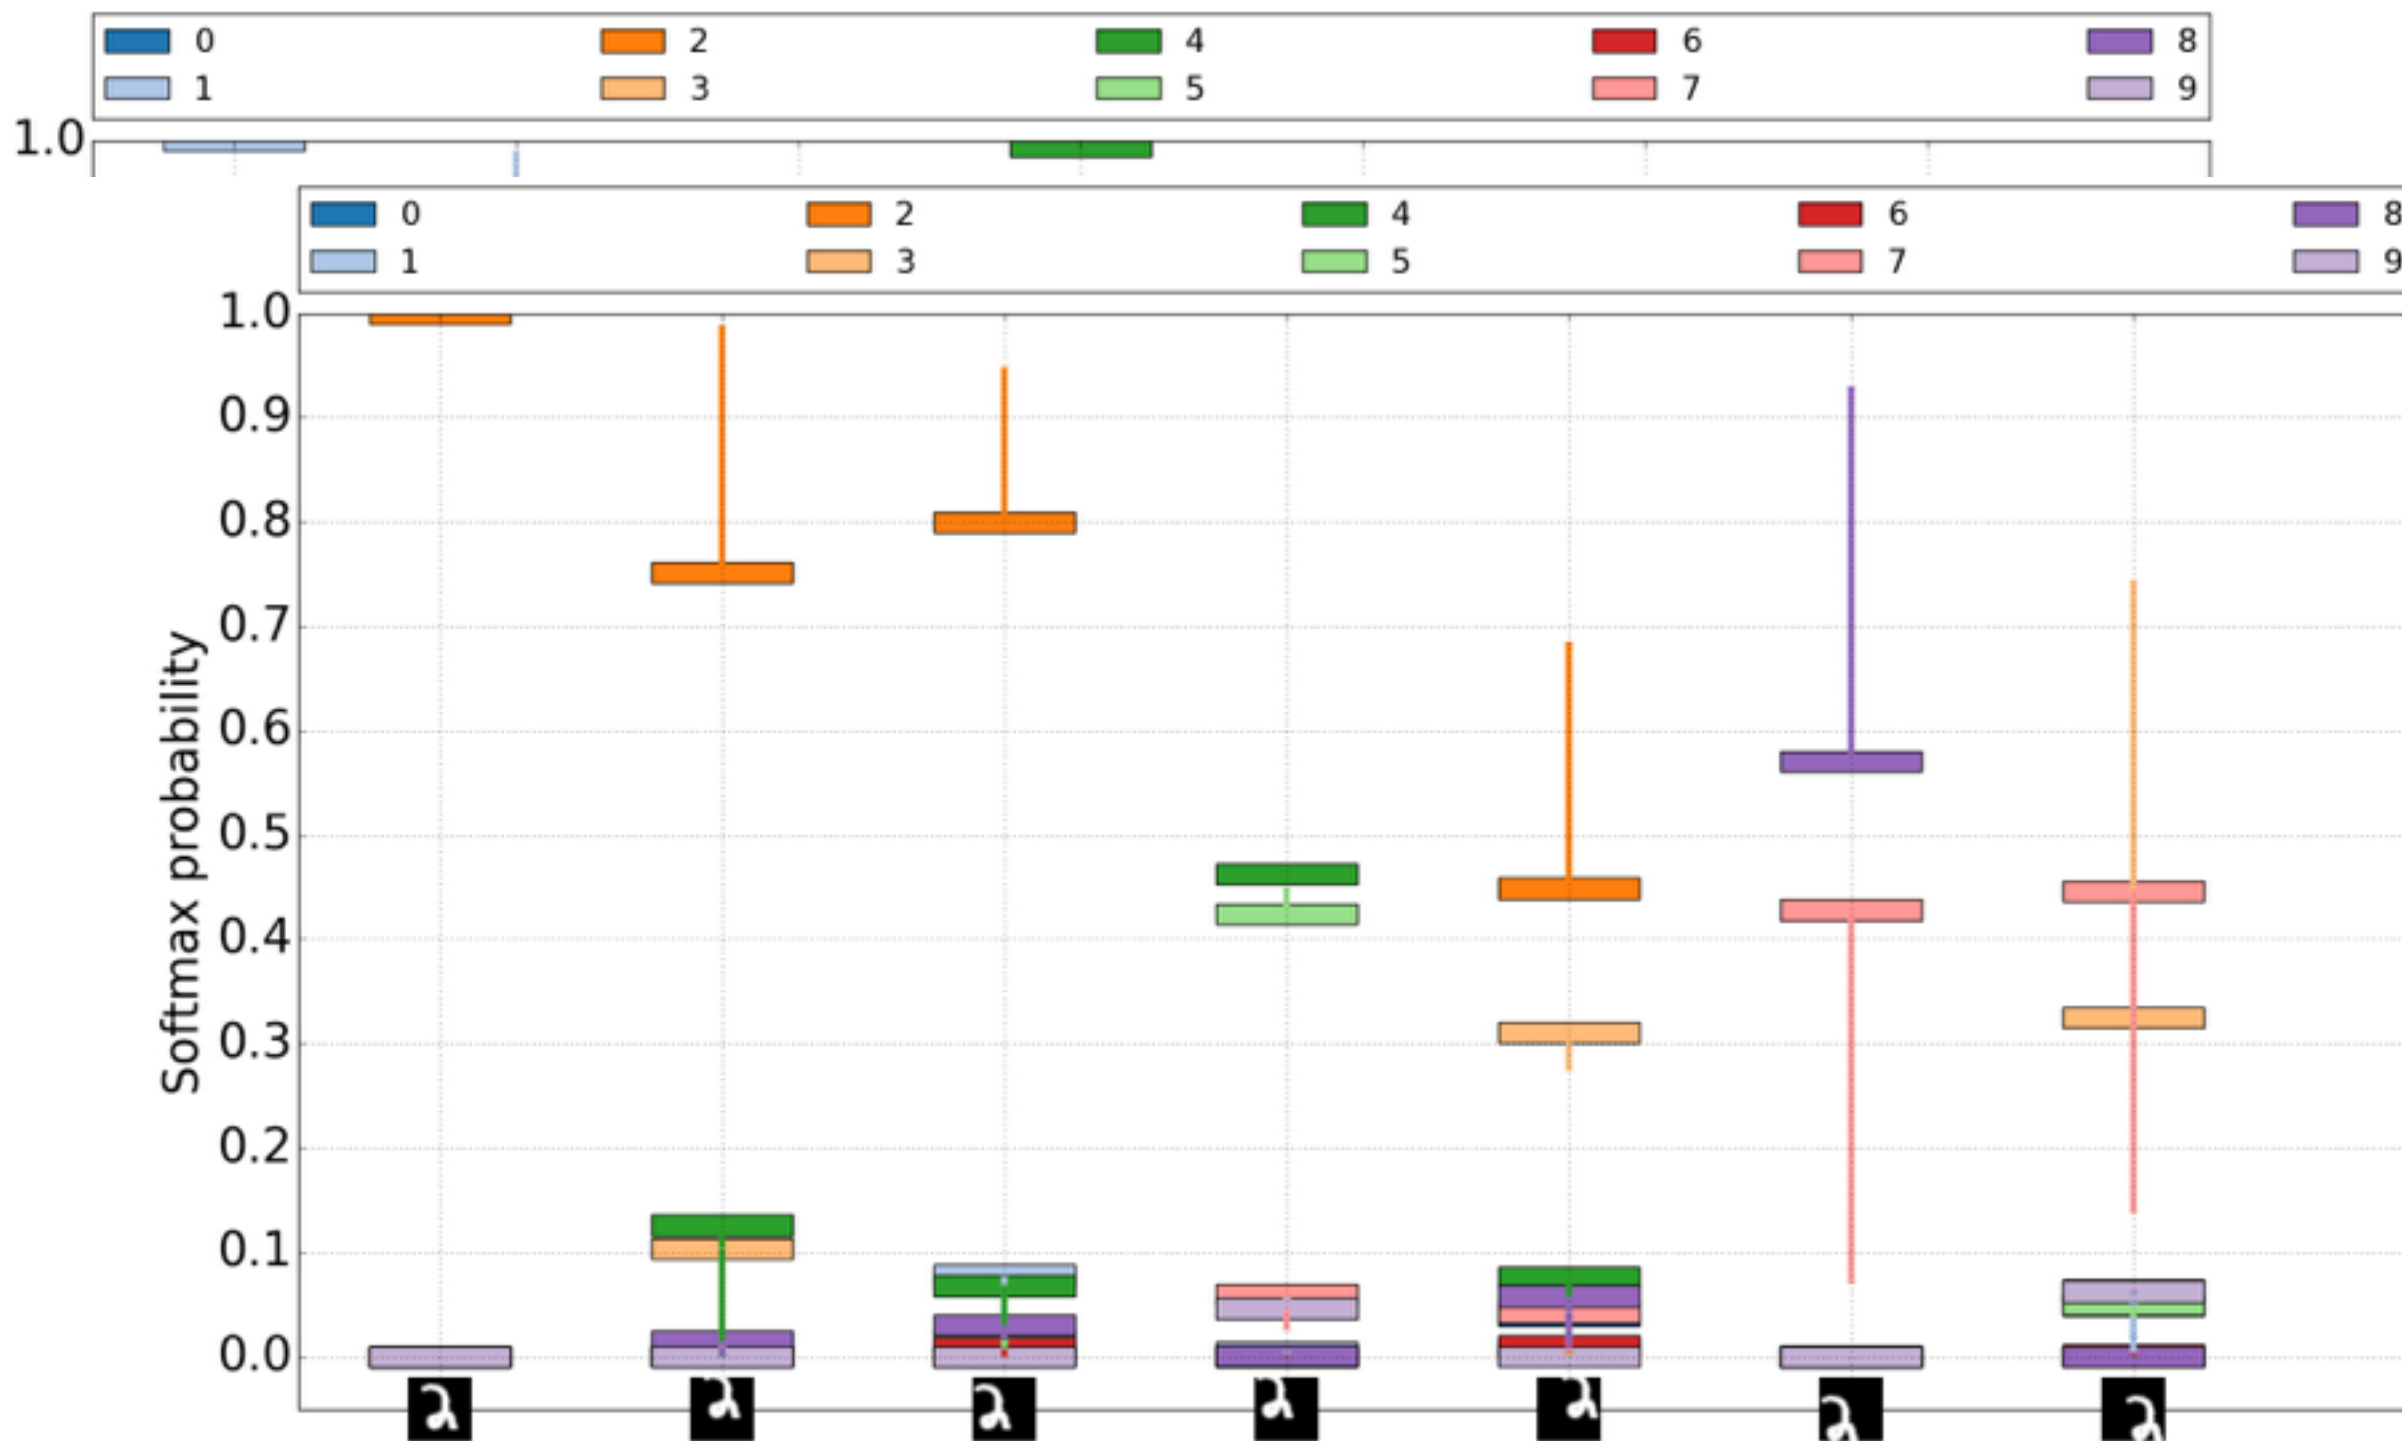
# SGLD translations

# SGLD translations






# SGLD translations

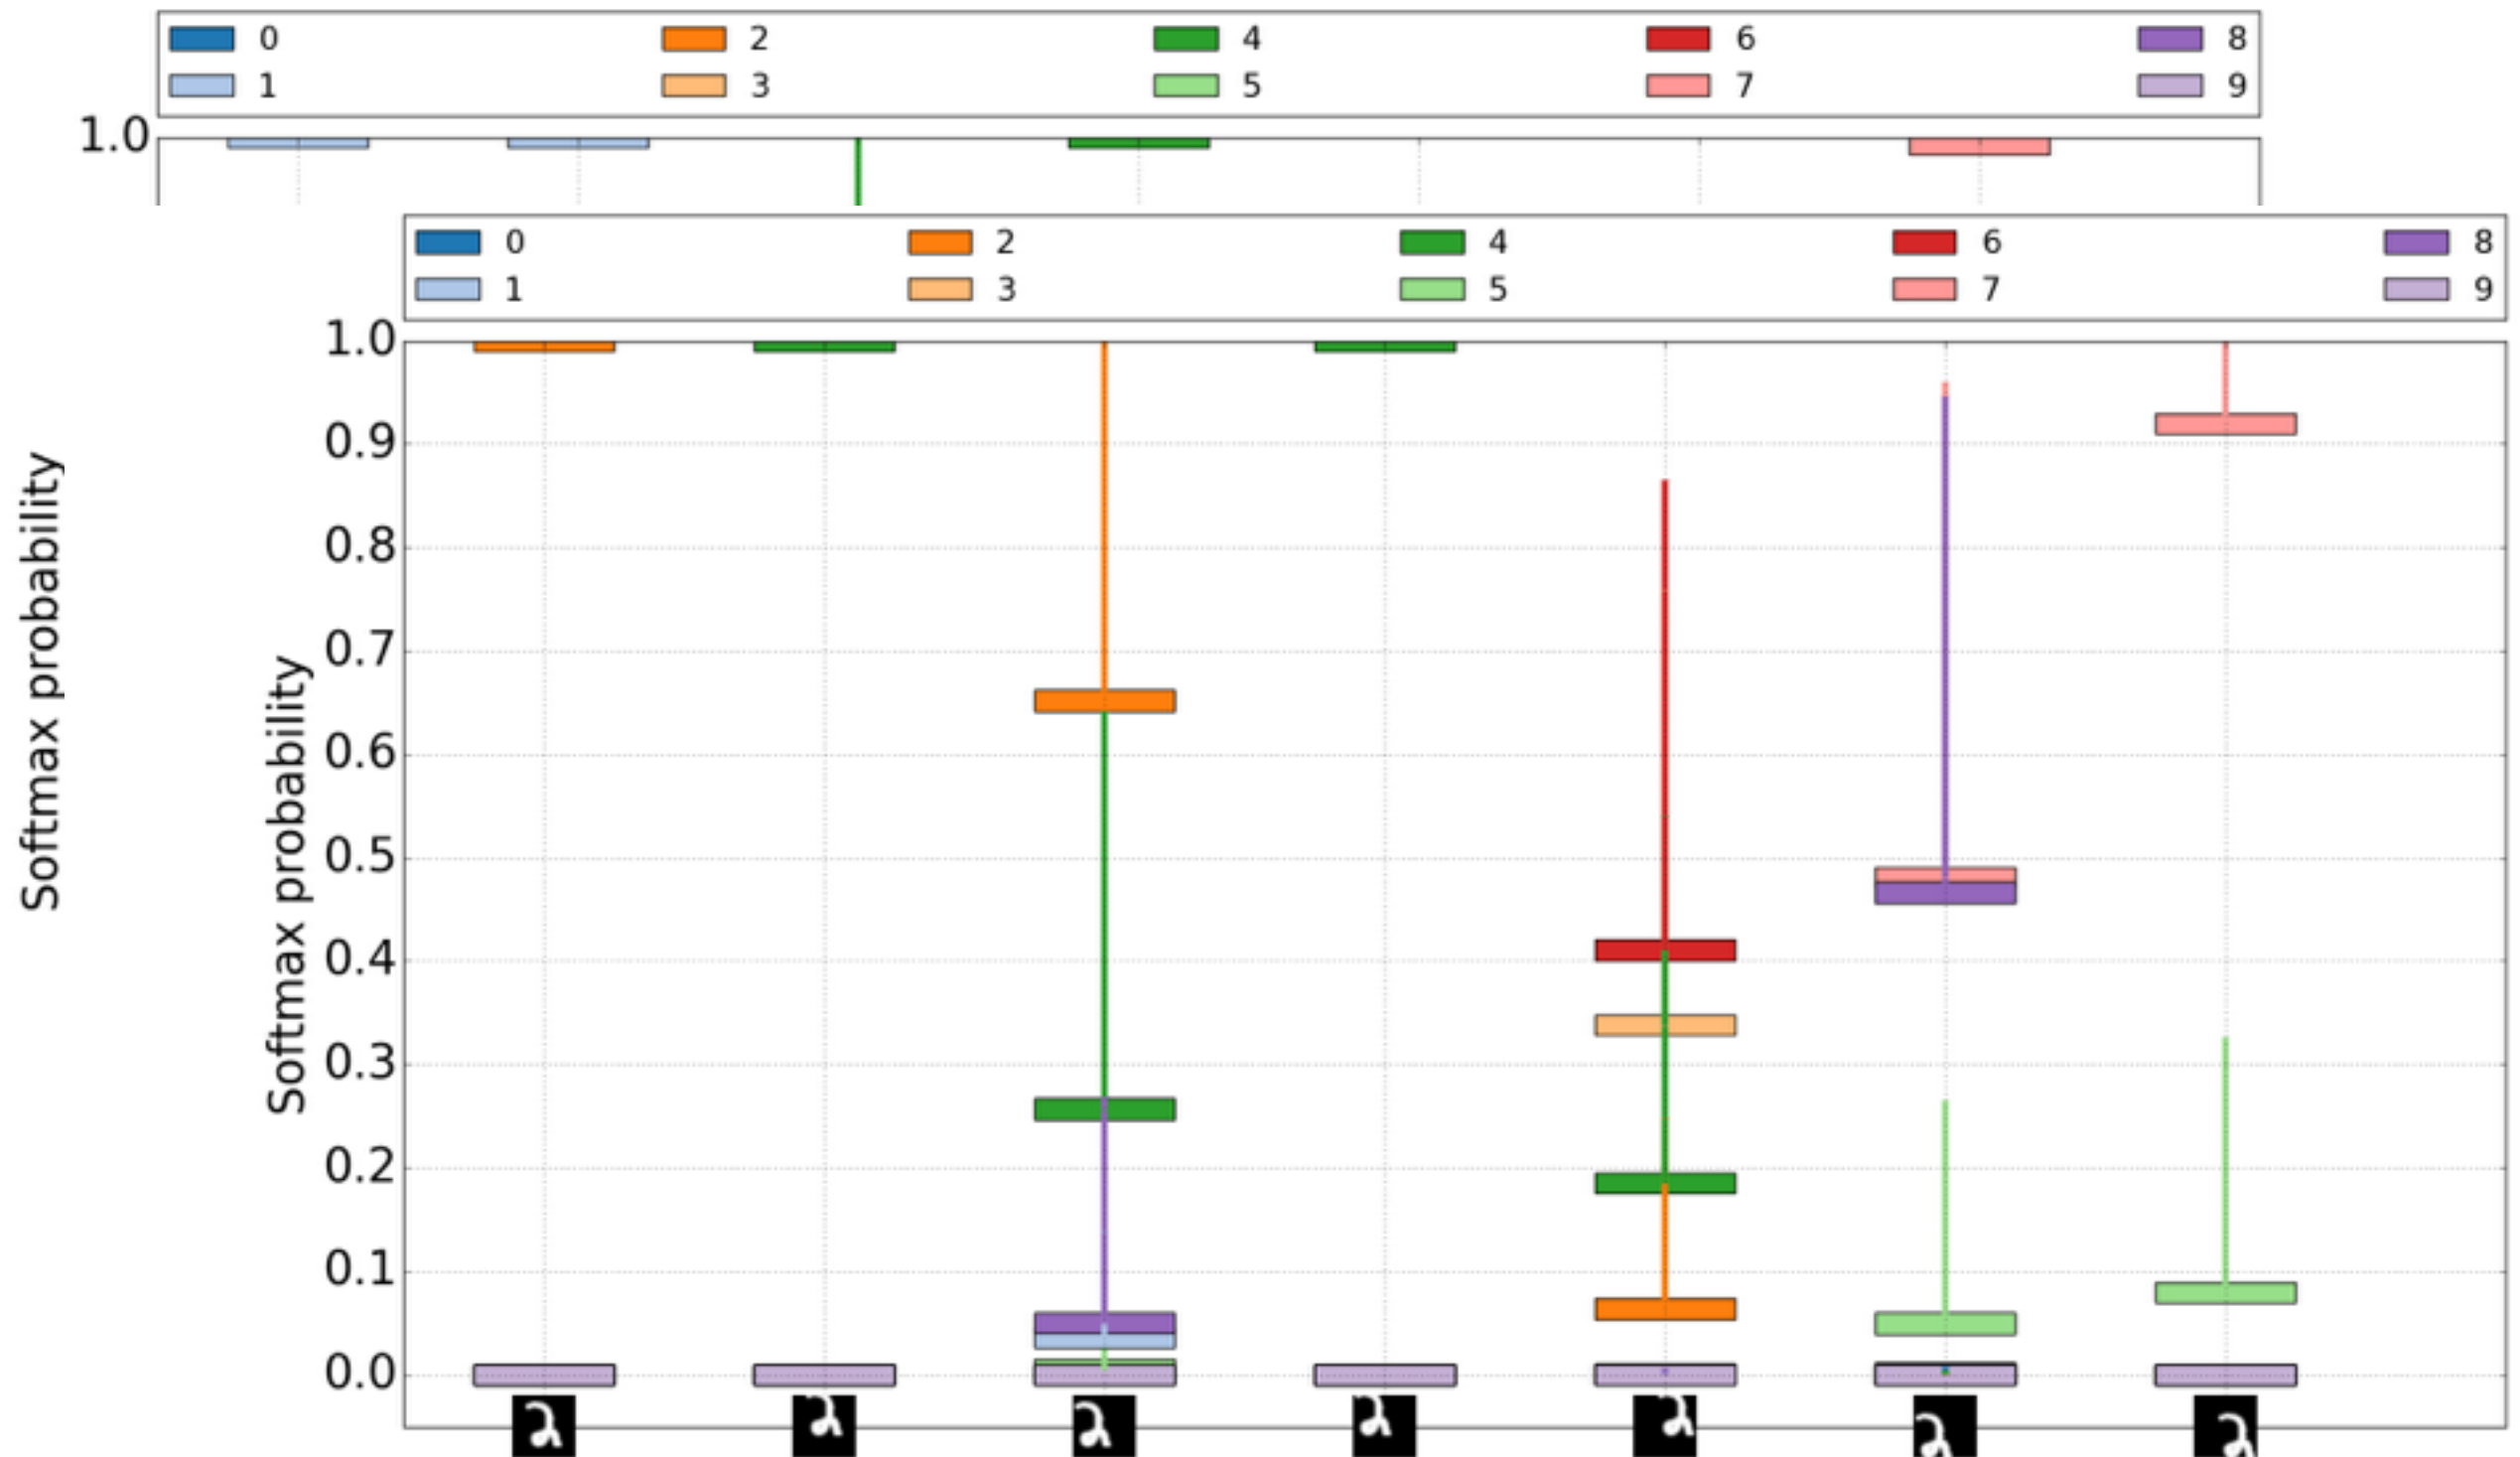


# Bootstrap translations

# Bootstrap translations




# Bootstrap translations



# Weight decay scaling

# Weight decay scaling




# Weight decay scaling




# Dropout scaling



# Dropout scaling




# Dropout scaling



# Matrix Gaussian scaling

# Matrix Gaussian scaling




# Matrix Gaussian scaling



# SGLD scaling

# SGLD scaling




# SGLD scaling





# Bootstrap scaling

# Bootstrap scaling



# Bootstrap scaling



# Verdict

# Verdict

- It seems that in all circumstances the uncertainties on inputs far from the data distribution are not good

# Verdict

- It seems that in all circumstances the uncertainties on inputs far from the data distribution are not good
- Deterministic weights are most affected and with Bayesian methods we do get away sometimes
  - But still in most cases we have erroneous overconfidence

# Verdict

- It seems that in all circumstances the uncertainties on inputs far from the data distribution are not good
- Deterministic weights are most affected and with Bayesian methods we do get away sometimes
  - But still in most cases we have erroneous overconfidence
- In general it appears to not be a parameter prior problem as it affects even the frequentist bootstrap
  - Although a better prior might fix it

# Verdict

- It seems that in all circumstances the uncertainties on inputs far from the data distribution are not good
- Deterministic weights are most affected and with Bayesian methods we do get away sometimes
  - But still in most cases we have erroneous overconfidence
- In general it appears to not be a parameter prior problem as it affects even the frequentist bootstrap
  - Although a better prior might fix it
- This suggests that the model itself (neural network) is not capable (at least as they are right now) to output reasonable probabilities



Thanks!